
Bayesian Hierarchical Reinforcement Learning

Abstract

We describe an approach to incorporating Bayesian priors in the MAXQ framework for hierarchical reinforcement learning (HRL). We define priors on the primitive environment model and on task pseudo-rewards. Since models for composite tasks can be complex, we use a mixed model-based/model-free learning approach to find an optimal hierarchical policy. We show empirically that (i) our approach results in improved convergence over non-Bayesian baselines, given sensible priors, (ii) task hierarchies and Bayesian priors can be complementary sources of information, and using both sources is better than either alone, (iii) taking advantage of the structural decomposition induced by the task hierarchy significantly reduces the computational cost of Bayesian reinforcement learning and (iv) in this framework, task pseudo-rewards can be learned instead of being manually specified, leading to automatic learning of hierarchically optimal rather than recursively optimal policies.

1. Introduction

Reinforcement learning (RL) is a well known framework that formalizes decision making in unknown, uncertain environments. RL agents learn policies that map environment states to available actions while optimizing some measure of long-term utility. While various algorithms have been developed for RL (?), and applied successfully to a variety of tasks (?), the standard RL setting suffers from at least two drawbacks. First, it is difficult to scale standard RL approaches to large state spaces (the well-known “curse of dimensionality”). Second, vanilla RL approaches do not incorporate prior knowledge about the environment and good policies.

Hierarchical reinforcement learning (HRL) (?) attempts to address the scaling problem by simplifying the overall decision making problem in different ways. For example, we could introduce macro-operators that can be used by the agent as part of its policy. These operators carry out sequences of primitive actions, so planning at the level of these operators may result in simpler policies (?). Another idea is to decompose the

task’s overall value function in a variety of ways, for example by defining task hierarchies (?) or partial programs with choice points (?). Here, the structure of the decomposition provides several benefits: first, for the “higher level” subtasks, policies are defined by calling “lower level” subtasks (which may themselves be quite complex); as a result policies for higher level subtasks may be expressed compactly. Second, a task hierarchy or partial program can impose constraints on the space of policies by encoding knowledge about the structure of good policies and thereby reduce the search space. Third, learning within subtasks allows *state abstraction*, that is, some state variables can be ignored because they do not affect the policy within that subtask. This also simplifies the learning problem.

While HRL attempts to address the scalability issue, it does not take into account probabilistic prior knowledge the agent may have about the task. For example, the agent may have some idea about where high/low utility states may be located and what their utilities may be, or some idea about the approximate shape of the value function or policy. *Bayesian reinforcement learning* addresses this issue by incorporating priors on models (?), value functions (??) or policies (?). Specifying good priors leads to many benefits, including initial good policies, directed exploration towards regions of uncertainty, and faster convergence to the optimal policy.

In our work, we propose an approach that incorporates Bayesian priors into hierarchical reinforcement learning. We use the MAXQ framework for HRL (?), that decomposes the overall task into subtasks so that value functions of the individual subtasks can be combined to recover the value function of the overall task. We extend this framework by incorporating priors on the primitive environment model and on task pseudo-rewards. In order to avoid building models for composite tasks (which can be very complex), we adopt a mixed model-based/model-free learning approach. We empirically evaluate our algorithm to understand the effect of the priors in addition to the task hierarchy. Our experiments indicate that: (i) taking advantage of probabilistic prior knowledge can lead to faster convergence, even for HRL, (ii) task hierarchies and Bayesian

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

priors can be complementary sources of information, and using both sources is better than either alone, (iii) taking advantage of the task hierarchy can reduce the computational cost of Bayesian RL, which generally tends to be very computation-intensive, and (iv) task pseudo-rewards can be learned instead of being manually specified, leading to automatic learning of hierarchically optimal rather than recursively optimal policies. In this way Bayesian RL and HRL are synergistic: Bayesian RL improves convergence of HRL and can learn hierarchy parameters, while HRL can reduce the cost of Bayesian RL.

Our work assumes the probabilistic priors to be given in advance and focuses on learning with them. Other work has addressed the issue of obtaining these priors. For example, one source of prior information is multi-task reinforcement learning (??), where an agent solves a sequence of RL tasks, and uses the solutions of previous tasks to build priors over models or policies for future tasks. We also assume the task hierarchy is given. Other work has explored learning MAXQ hierarchies in different settings (?).

2. Background and Related Work

In the MAXQ framework, each composite subtask T_i defines a semi-Markov decision process with parameters $\langle S_i, X_i, C_i, G_i \rangle$. S_i is a predicate that defines the set of “non-terminal” states for T_i , where T_i may be called by its parent. G_i is a predicate that defines a set of “goal” states for T_i . The actions available within T_i are described by the set of “child tasks” C_i . Finally, X_i denotes the set of “relevant state variables” for T_i . Often, we unify S_i and (the negation of) G_i into a single “termination” predicate, P_i . A (state, action, next-state) (s, a, s') triple where $P_i(s)$ is false, $P_i(s')$ is true, $a \in C_i$, and the transition probability $P(s'|s, a) > 0$ is called an *exit* of the subtask T_i . A pseudo-reward function $\hat{R}(s, a)$ can be defined over exits to express preferences over the possible exits of a subtask.

A subtask can be invoked in any state s where $P_i(s)$ is false, and it terminates when an exit causes $P_i(s')$ to become true. The set S_i is defined using a projection function that maps a world state to an abstract state defined by a subset of the state variables. A hierarchical policy π for the overall task is an assignment of a local policy to each SMDP T_i . A *hierarchically optimal policy* is a hierarchical policy that has the maximum expected reward. A hierarchical policy is said to be *recursively optimal* if the local policy for each subtask is optimal given that all its subtask policies are optimal. Recursively optimal policies can be obtained by fixing all pseudo-rewards to zero. Given a task

graph, model-free (?) or model-based (?) methods can be used to learn value functions for each task-subtask pair. In the model-free method, a policy is produced by maintaining a value and a *completion* function for each subtask. For a task i , the value $V(a, s)$ denotes the expected value of calling child task a in state s . This is (recursively) estimated as the expected reward obtained while executing a . The completion function $C(i, s, a)$ denotes the expected reward obtained while *completing* i after having called a in s . The central idea behind MAXQ is that the value of i , $V(i, s)$, can be (recursively) decomposed in terms of $V(a, s)$ and $C(i, s, a)$.

Bayesian reinforcement learning methods incorporate probabilistic prior knowledge on models (?), value functions (??), policies (?) or combinations (?). One Bayesian model-based RL algorithm proceeds as follows. At each step, a distribution over model parameters is maintained. At each step, a model is sampled from this distribution (Thompson sampling (??)). This model is then solved and actions are taken according to the policy obtained. This yields observations about the environment model. These are used to update the parameters of the current distribution to create a posterior distribution over models. This procedure is then iterated to convergence. Variations of this idea have been investigated; for example, some work converts the distribution over models to an empirical distribution over Q -functions, and produces policies by sampling from this distribution instead (?).

Relatively little work exists that attempts to incorporate probabilistic priors into HRL. We have found one preliminary attempt (?). This work builds on the RMAX+MAXQ (?) method, which extends RMAX (?) to MAXQ by maintaining models for each subtask and performing RMAX-style exploration in each. The Bayesian approach adds priors onto each subtask model and performs (separate) Bayesian model-based learning for each subtask.¹ In our approach, we do not construct models for subtasks, which can be very complex in general. Instead, we only maintain distributions over primitive actions, and use a mixed model-based/model-free learning algorithm that naturally integrates into the standard MAXQ learning algorithm.

3. Bayesian MAXQ Algorithm

In this section, we describe our approach to incorporating probabilistic priors into MAXQ. We use priors over

¹While we believe this description is accurate, unfortunately, due to language issues and some missing technical and experimental details in the cited article, we have been unable to replicate this work.

two quantities: primitive models and pseudo-rewards. As we explain below, pseudo-rewards are value functions; thus our approach uses priors both on models and value functions. While such an integration may not be needed for standard Bayesian RL, it appears naturally in our setting.

We first describe our approach to incorporating priors on environment models alone (assuming pseudo-rewards are fixed). We do this following the Bayesian model-based RL framework. At each step we have a distribution over environment models (initially the prior). The algorithm has two main subroutines: the main BAYESIAN_MAXQ routine (Algorithm ??) and an auxiliary RECOMPUTE_VALUE routine (Algorithm ??). In this description, the value V and completion C functions are assumed to be global.

At the start of each episode, the BAYESIAN_MAXQ routine is called with the *Root* task and the initial state for the current episode. The MAXQ execution protocol is then followed, where each task chooses an action based on its current value function (initially random). When a primitive action is reached and executed, it updates the posterior over model parameters (Line ??) and its own value estimate (which is just the reward function for primitive actions). When a task exits and returns to its parent, the parent subsequently updates its completion function based on the current estimates of the value of the exit state (Lines ?? and ??). Note that in MAXQ, the value function of a composite task can be (recursively) computed using the completion functions of subtasks and the rewards obtained by executing primitive actions, so we do not need to separately store or update the value functions (except for the primitive actions where the value function is the reward). Finally, each primitive action maintains a count of how many times it has been executed and each composite task maintains a count of how many child actions have been taken.

When k (an algorithm parameter) steps have been executed in a composite task, BAYESIAN_MAXQ calls RECOMPUTE_VALUE to re-estimate the value and completion functions (the check on k is shown in RECOMPUTE_VALUE, Line ??). When activated, this function recursively re-estimates the value/completion functions for all subtasks of the current task. At the level of a primitive action, this simply involves resampling the reward and transition parameters from the current posterior over models. For a composite task, we use the MAXQ-Q algorithm (Table 4 in (?)). We run this algorithm for Sim episodes, starting with the current subtask as the root, with the current pseudo-reward estimates (we explain below how these are obtained).

Algorithm 1 BAYESIAN_MAXQ

Input: Task i , State s , Update Interval k , Simulation Episodes Sim
Output: Next state s' , steps taken N , cumulative reward CR

```

1: if  $i$  is primitive then
2:   Execute  $i$ , observe  $r$ ,  $s'$ 
3:   Update current posterior parameters  $\Psi$  using  $(s, i, r, s')$ 
4:   Update current value estimate:  $V(i, s) \leftarrow (1 - \alpha) \cdot V(i, s) + \alpha \cdot r$ 
5:    $Count(i) \leftarrow Count(i) + 1$ 
6:   return  $(s', 1, r)$ 
7: else
8:    $N \leftarrow 0$  { $i$  is composite}
9:    $CR \leftarrow 0$ 
10:   $taskStack \leftarrow Stack()$ 
11:  while  $i$  is not terminated do
12:    RECOMPUTE_VALUE( $i, k, Sim$ )
13:     $a \leftarrow \epsilon$ -greedy action from  $V(i, s)$ 
14:     $\langle s', N_a, cr \rangle \leftarrow \text{BAYESIAN\_MAXQ}(a, s)$ 
15:     $taskStack.push(\langle a, s', N_a, cr \rangle)$ 
16:     $a_{s'}^* \leftarrow \arg \max_{a'} [\tilde{C}(i, s', a') + V(a', s')]$ 
17:     $C(i, s, a) \leftarrow (1 - \alpha) \cdot C(i, s, a) + \alpha \cdot \gamma^{N_a} [C(i, s', a_{s'}^*) + V(a_{s'}^*, s')]$ 
18:     $\tilde{C}(i, s, a) \leftarrow (1 - \alpha) \cdot \tilde{C}(i, s, a) + \alpha \cdot \gamma^{N_a} [\tilde{R}(i, s') + \tilde{C}(i, s', a_{s'}^*) + V(a_{s'}^*, s')]$ 
19:     $s \leftarrow s'$ 
20:     $CR \leftarrow CR + \gamma^N \cdot cr$ 
21:     $N \leftarrow N + N_a$ 
22:     $Count(i) \leftarrow Count(i) + 1$ 
23:  end while
24:  UPDATE_PSEUDO_REWARD( $taskStack, \tilde{R}(i, s')$ )
25:  return  $(s', N, CR)$ 
26: end if
    
```

This algorithm recursively updates the completion function of the task graph below the current task. Note that in this step, the subtasks with primitive actions use model-based updates. That is, when a primitive action is “executed” in such tasks, the currently sampled transition function (part of Θ in Line ??) is used to find the next state, and then the associated reward is used to update the completion function. This is similar to Lines ??, ?? and ?? in BAYESIAN_MAXQ, except that it uses the sampled model Θ instead of the real environment. After RECOMPUTE_VALUE terminates, a new set of value/completion functions are available for BAYESIAN_MAXQ to use to select actions.

Next we discuss the issue of task pseudo-rewards. A pseudo-reward is a value associated with a subtask exit that defines how “good” that exit is for that subtask.

Algorithm 2 RECOMPUTE_VALUE

Input: Task i , Update Interval k , Simulation Episodes Sim

Output: Recomputed value and completion functions for the task graph below and including i

- 1: **if** $Count(i) < k$ **then**
- 2: **return**
- 3: **end if**
- 4: **if** i is primitive **then**
- 5: Sample new transition and reward parameters Θ from current posterior Ψ
- 6: **else**
- 7: **for all** child tasks a of i **do**
- 8: RECOMPUTE_VALUE(a, k, Sim)
- 9: **end for**
- 10: **for** Sim episodes **do**
- 11: $s \leftarrow$ random nonterminal state of i
- 12: Run MAXQ-Q(i, s, Θ, \tilde{R})
- 13: **end for**
- 14: **end if**
- 15: $Count(i) \leftarrow 0$

Algorithm 3 UPDATE_PSEUDO_REWARD

Input: $taskStack$, Parent’s pseudo reward \tilde{R}_p

- 1: $tempCR \leftarrow \tilde{R}_p$
- 2: $N_{a'} \leftarrow 0, cr' \leftarrow 0$
- 3: **while** $taskStack$ is not empty **do**
- 4: $\langle a, s, N_a, cr \rangle \leftarrow taskStack.pop()$
- 5: $tempCR \leftarrow \gamma^{N_a'} \cdot tempCR + cr'$
- 6: Update pseudo-reward posterior Φ for $\tilde{R}(a, s)$ using $(a, s, tempCR)$
- 7: Resample $\tilde{R}(a, s)$ from Φ
- 8: $N_{a'} \leftarrow N_a, cr' \leftarrow cr$
- 9: **end while**

The *ideal* pseudo-reward for an exit is the expected reward under the hierarchically optimal policy after exiting the subtask, until the global task (Root) ends; thus the pseudo-reward is really a value function. This pseudo-reward would enable the subtask to choose the “right” exit *in the context of* what the rest of the task hierarchy is doing. In standard MAXQ, these have to be set manually. This is problematic because it presupposes (quite detailed) knowledge of the hierarchically optimal policy. Further, setting the *wrong* pseudo-rewards can result in non-convergence or highly suboptimal policies. Sometimes this problem is sidestepped simply by setting all pseudo-rewards to zero, resulting in recursively optimal policies. However, it is easy to construct examples where a recursively optimal policy is arbitrarily worse than the hierarchically optimal policy. For all these reasons, pseudo-rewards are major

“nuisance parameters” in the MAXQ framework.

What makes learning pseudo-rewards tricky is that they are not only value functions, but also function as *parameters* of MAXQ. That is, setting different pseudo-rewards essentially results in a new learning problem. For this reason, simply trying to learn pseudo-rewards in a standard temporal difference (TD) way fails (as we show in our experiments). In our work, we take advantage of the fact that Bayesian RL allows us to address both these issues. First, we can treat value functions as probabilistic unknown parameters. Second, and more importantly, a key idea in Bayesian RL is the “lifting” of exploration to the space of task parameters. That is, instead of exploration through action selection, Bayesian RL can perform exploration by sampling task parameters. Thus treating a pseudo-reward as an unknown Bayesian parameter also leads to *exploration over the value of this parameter*, until an optimal value is found. In this way, hierarchically optimal policies can be learned from scratch—a major advantage over the standard MAXQ setting.

To learn pseudo-rewards, we again maintain a distribution over all such parameters, Φ , initially a prior. For simplicity, we only focus on tasks with multiple exits, since otherwise, a pseudo-reward has no real effect. When a composite task executes, we keep track of each child task’s execution in a stack. When the parent itself exits, we obtain a new observation of the pseudo-rewards of each child by computing the discounted cumulative reward received *after* it exited, added to the current estimate of the parent’s pseudo-reward (Algorithm ??). This observation is used to update the current posterior over the child’s pseudo-reward. Since this is a value function estimate, early in the learning process, the estimates are noisy. Following prior work (?), we use a window containing the most recent observations. When a new observation arrives, the oldest observation is removed, the new one is added and a new posterior estimate is computed. After updating the posterior, it is sampled to obtain a new pseudo-reward estimate for the associated exit. This estimate is used where needed (in Algorithms ?? and ??) until the next posterior update. Combined with the model-based priors above, we hypothesize that this procedure, iterated till convergence, will produce a hierarchically optimal policy.

4. Empirical Evaluation

In this section, we evaluate our approach and test four hypotheses: First, does incorporating model-based priors help speed up the convergence of MAXQ to the optimal policy? Second, how does the task hierarchy

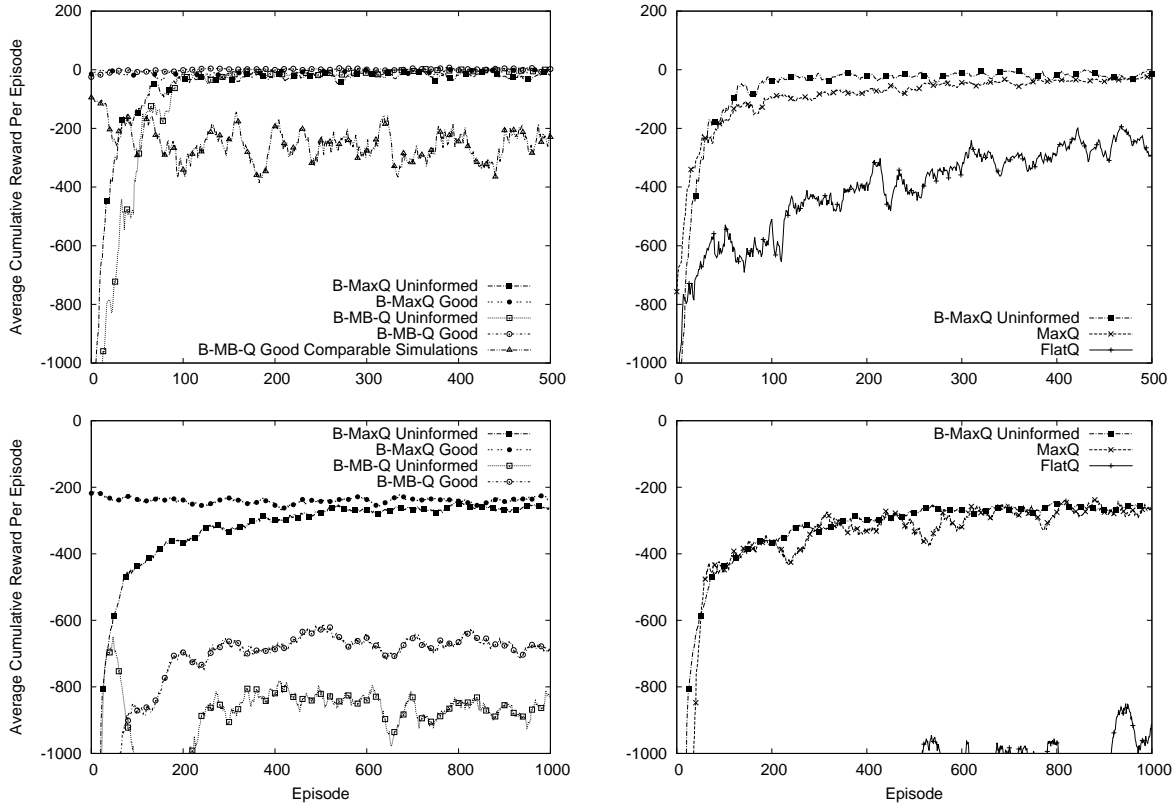


Figure 1. Performance of different algorithms on Taxi-World (top row) and Resource-collection (bottom row). The prefix “B-” denotes Bayesian, “Uninformed/Good” denotes the prior and “MB” denotes model-based.

interact with the prior? Does the task hierarchy still matter if very good priors are available for primitive actions? Third, how does Bayesian MAXQ compare to standard (flat) Bayesian RL? Does Bayesian RL perform better (in terms of computational time) if a task hierarchy is available? Finally, can our approach effectively learn pseudo-rewards and policies that are hierarchically optimal?

We first focus on evaluating the first three hypotheses using domains where a zero pseudo-reward results in hierarchical optimality. To evaluate these hypotheses, we use two domains: the fickle version of Taxi-World (?) (625 states) and Resource-collection (?) (8265 states). In Taxi-World, the agent controls a taxi in a grid-world that has to pick up a passenger from a source location and drop them off at their destination. The state variables consist of the location of the taxi and the source and destination of the passenger. The actions available to the agent consist of navigation actions and actions to pickup and putdown the passenger. The agent gets a reward of +20 upon completing the task, a constant -1 reward for every action and a -10 penalty for an erroneous action. Further, each navigation action has a 15% chance of moving in a direction orthogonal to the intended move.

The Resource-collection domain is inspired by real-time strategy games where one component involves collecting resources from a map. We simulate this through a grid world environment where the agent controls a unit that can harvest gold and chop wood. Here the state variables consist of the location of the agent, what the agent is carrying, whether a goldmine or forest is adjacent to its current location and whether a desired gold or wood quota has been met. The actions available to the agent are to move to a specific location, chop gold or harvest wood, and to deposit the item it is carrying (if any). As before, for each navigation action the agent has a probability of moving to a random location with 0.3 probability. In our experiments, the map contains two goldmines and two forests, each containing two units of gold and two units of wood, and the gold and wood quota is set to three each. The agent gets a +50 reward when it meets the gold/wood quota, a constant -1 reward for every action and an additional -1 for erroneous actions (such as trying to deposit when it is not carrying anything).

For the Bayesian methods, we use Dirichlet priors for the transition function parameters and Normal-Gamma priors for the reward function parameters. We use two priors: an uninformed prior, set to ap-

Table 1. CPU time for methods on Taxi-World.

Method	Time for 500 Episodes (s)
Bayesian MaxQ, Uninformed Prior	179
Bayesian Model-based Q, Uninformed Prior	232
Bayesian MaxQ, Good Prior	14
Bayesian Model-based Q, Good Prior	119
Bayesian Model-based Q, Good Prior & Comparable Simulations	934
MaxQ	1.15
FlatQ	0.53

proximate a uniform distribution, and a “good” prior where a previously computed model posterior is used as the “prior.” The prior distributions we choose are conjugate to the likelihood, so we can compute the posterior distributions over the relevant parameters in closed form. In general, this is not necessary; more complex priors could be used as long as we can sample from the posterior distribution.

The methods we use in this experiment are: (i) Flat Q, the standard, non-Bayesian Q-learning algorithm, (ii) MAXQ-0, the standard, non-Bayesian Q-learning algorithm for MAXQ with no pseudo-reward, (iii) Bayesian model-based Q-learning with an uninformed prior and (iv) a “good” prior, (v) Bayesian MAXQ (our proposed approach) with an uninformed prior and (vi) with a “good” prior. In our implementation, the Bayesian model-based Q-learning uses the same code as the Bayesian MAXQ algorithm, with a “trivial” hierarchy consisting of the Root task with only the primitive actions as children. For the Bayesian methods, the update frequency k was set to 50 for Taxi-World and 25 for Resource-collection. Sim was set to 200 for Bayesian MAXQ for Taxi-World and 1000 for Bayesian model-based Q, and to 1000 for both for Resource collection. The results are shown in Figure ??.

From these results, comparing the Bayesian versions of MAXQ to standard MAXQ, we observe that for Taxi-World, the Bayesian version converges faster to the optimal policy even with the uninformed prior, while for Resource-collection, the convergence rates are similar. When a good prior is available, convergence is very fast (almost immediate) in both domains. Thus, the availability of model priors can help speed up convergence in many cases for HRL.

Next, we compare the Bayesian MAXQ approach to “flat” Bayesian model-based Q learning. We note that in Taxi-World, with uninformed priors, though the “flat” method initially does worse, it soon catches

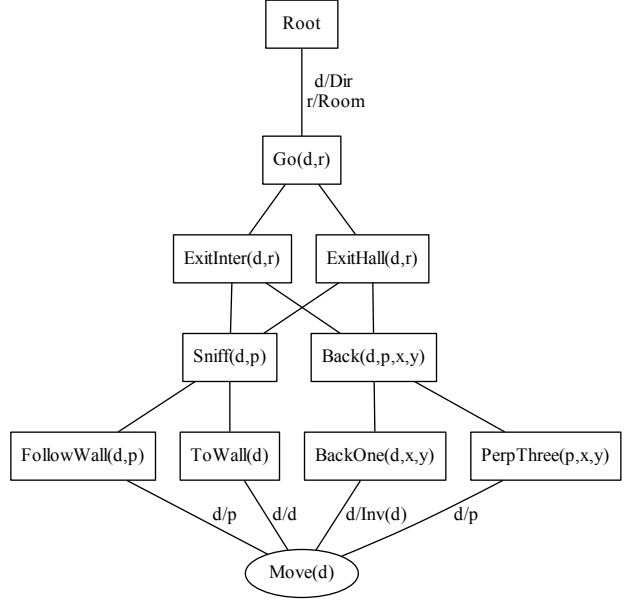


Figure 2. Task Hierarchy for Hallway.

up to standard MAXQ and then to Bayesian MAXQ. This is probably because in this domain, the primitive models are relatively easy to acquire, and the task hierarchy provides no additional leverage. For Resource-collection, however, even with a good prior, “flat” Bayesian model-based Q does not converge. The difference is that in this case, the task hierarchy encodes extra information that cannot be deduced just from the models. In particular, the task hierarchy tells the agent that good policies consist of gold/wood collection moves followed by deposit moves. Since the reward structure in this domain is very sparse, it is difficult to deduce this even if very good models are available. Taken together, these results show that task hierarchies and model priors can be complementary: in general, Bayesian MAXQ outperforms both flat Bayesian RL and MAXQ (in speed of convergence, since here MAXQ can learn the optimal policy).

Next, we compare the time taken by the different approaches in our experiments in Taxi-World (Table ??). As expected, the Bayesian RL approaches are significantly slower than the non-Bayesian approaches. However, out of the Bayesian methods, the Bayesian MAXQ approaches are significantly faster than the flat Bayesian model-based approaches. This can be attributed to the fact that for the flat case, during the simulation in RECOMPUTE_VALUE, a much larger task needs to be solved, while the Bayesian MAXQ approach is able to take into account the structure of the hierarchy to only simulate subtasks as needed, which ends up being much more efficient. However, we note that we

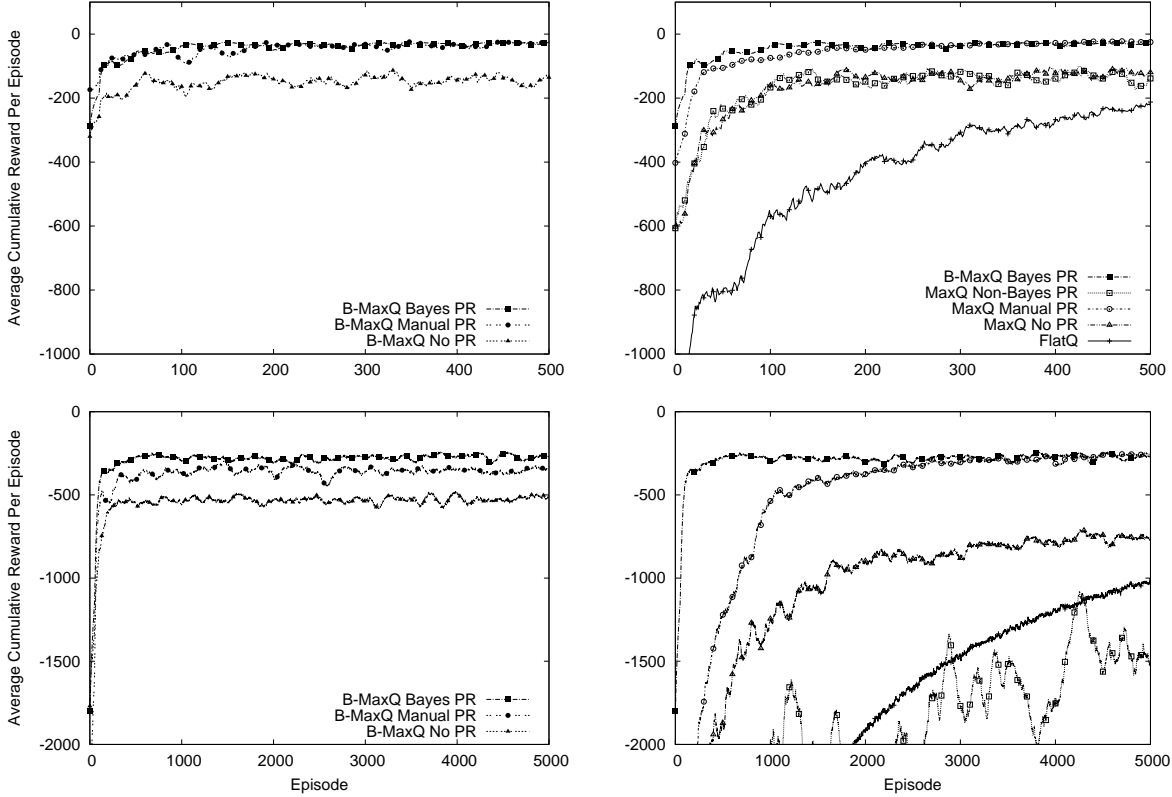


Figure 3. Performance of different algorithms on Modified-Taxi-World (top row) and Hallway (bottom row). The prefix “B-” denotes Bayesian, “PR” denotes Pseudo Reward.

allowed the flat Bayesian model-based approach 1000 episodes of simulation as opposed to 200 for Bayesian MAXQ. Clearly this increases the time taken for the flat cases. But at the same time, this is necessary: the “Comparable Simulations” row (and curve in Figure ?? top left) shows that, if the simulations are reduced to 250 episodes for this approach, the resulting values are no longer reliable and the performance of the Bayesian flat approach drops sharply. Thus, taking advantage of the hierarchical task decomposition helps reduce the computational cost of Bayesian RL.

Finally we evaluate how well our approach estimates pseudo-rewards. Here we use two domains: a Modified-Taxi-World and a Hallway domain (??) (4320 states). In Modified-Taxi-World, We allow dropoffs at any one of the four specific locations and do not provide a reward for task termination. Thus the Navigate subtask needs a pseudo-reward (corresponding to the correct dropoff location) to learn a good policy. The Hallway domain consists of a maze with a large scale structure of hallways and intersections. The agent has stochastic movement actions. The task hierarchy is shown in Figure ??.

For these experiments, we use uninformed priors on

the environment model. The pseudo-reward priors are set to prefer each exit from a subtask equally. The baselines we use are: (i) Bayesian MAXQ and MAXQ with fixed zero pseudo-reward, (ii) Bayesian MAXQ and MAXQ with fixed manually set pseudo-reward, (iii) flat Q and (iv) MAXQ with a non-Bayesian pseudo-reward update. This last method tracks pseudo-reward just as our approach; however, instead of a Bayesian update, it updates the pseudo-reward using a temporal difference update, treating it as a simple value function. The results are shown in Figure ??.

From these results, we first observe that the methods with zero pseudo-reward always do worse than those with “proper” pseudo-reward, indicating that in these cases the recursively optimal policy is not the hierarchically optimal policy. When a pseudo-reward is manually set, in both domain, MAXQ converges to better policies. We observe that in each case, the Bayesian MAXQ approach is able to learn a policy that is as good, starting with no pseudo rewards; further, its convergence rates are often better. Further, we observe that the simple TD update strategy (MaxQ Non-Bayes PR in Figure ??) fails in both cases—in Modified-Taxi-World, it is able to learn a policy that is approximately as good as a recursively optimal policy, but in the Hall-

way domain, it fails to converge completely, indicating that this strategy cannot generally learn pseudo-rewards. These results indicate that a strategy that incorporates Bayesian methods into MAXQ can successfully learn pseudo-rewards from scratch and produce hierarchically optimal policies.

5. Conclusion

In this paper, we have proposed an approach to incorporating probabilistic priors on environment models and task pseudo-rewards into HRL by extending the MAXQ framework. Our approach uses a combination of model-based and model-free learning to compute a hierarchically optimal policy. Our experiments indicate that several synergies exist between HRL and Bayesian RL, and combining them is fruitful. In future work, we plan to investigate more complex domains, where models and value functions cannot be exactly represented, as well as investigate multi-task RL scenarios where the priors we currently assume to be given can be learned.

References

- Andre, D. and Russell, S. State Abstraction for Programmable Reinforcement Learning Agents. In *Proceedings of AAAI*, 2002.
- Barto, Andrew G. and Mahadevan, Sridhar. Recent advances in hierarchical reinforcement learning. *Discrete Event Dynamic Systems*, 13(4):341–379, 2003.
- Brafman, Ronen I., Tennenholtz, Moshe, and Kaelbling, Pack. R-max - a general polynomial time algorithm for near-optimal reinforcement learning. *Journal of Machine Learning Research*, 2001.
- Dai, Zhaohui, Chen, Xin, Cao, Weihua, and Wu, Min. Model-based learning with bayesian and maxq value function decomposition for hierarchical task. In *Proceedings of the 8th World Congress on Intelligent Control and Automation*, 2010.
- Dearden, R., Friedman, N., and Russell, S. Bayesian Q-learning. In *Proceedings of the Fifteenth National Conference on Artificial Intelligence*, 1998. URL <http://www.cs.huji.ac.il/~nir/Papers/DFR1.ps>.
- Dearden, R., Friedman, N., and Andre, D. Model based bayesian exploration. In *Proceedings of Fifteenth Conference on Uncertainty in Artificial Intelligence*. Morgan Kaufmann, 1999.
- Dietterich, Thomas G. Hierarchical reinforcement learning with the maxq value function decomposition. *Journal of Artificial Intelligence Research*, 13:227–303, 2000.
- Engel, Y., Mannor, S., and Meir, R. Bayes meets Bellman: the Gaussian process approach to temporal difference learning. In *Proceedings of the 20th International Conference on Machine Learning*, 2003. URL www.cs.ualberta.ca/~yaki/papers/icml_gptd.ps.
- Ghavamzadeh, M. and Engel, Y. Bayesian actor-critic algorithms. In Ghahramani, Zoubin (ed.), *Proceedings of the 24th Annual International Conference on Machine Learning (ICML 2007)*, pp. 297–304. Omnipress, 2007a.
- Ghavamzadeh, Mohammad and Engel, Yaakov. Bayesian policy gradient algorithms. In *Advances in Neural Information Processing Systems 19*. MIT Press, 2007b.
- Jong, Nicholas K. and Stone, Peter. Hierarchical model-based reinforcement learning: R-max + maxq. In *Proceedings of the 25th International Conference on Machine Learning*, 2008.
- Kaelbling, Leslie Pack, Littman, Michael L., and Moore, Andrew W. Reinforcement learning: A survey. *Journal of Artificial Intelligence Research*, 4:237–285, 1996.
- Lazaric, Alessandro and Ghavamzadeh, Mohammad. Bayesian multi-task reinforcement learning. In *Proc. 27th International Conference on Machine Learning*, 2010.
- Mehta, N., Ray, S., Tadepalli, P., and Dietterich, T. Automatic discovery and transfer of MAXQ hierarchies. In McCallum, Andrew and Roweis, Sam (eds.), *Proceedings of the 25th International Conference on Machine Learning*, pp. 648–655. Omnipress, 2008.
- Parr, Ronald Edward. *Hierarchical Control and Learning for Markov Decision Processes*. PhD thesis, 1998.
- Stolle, Martin and Precup, Doina. *Learning Options in reinforcement Learning*, volume 2371/2002 of *Lecture Notes in Computer Science*, pp. 212–223. Springer, 2002.
- Strens, M. J. A. A Bayesian framework for reinforcement learning. In *Proceeding of the 17th International Conference on Machine Learning*, 2000. URL <http://uk.geocities.com/mjstrens/icml.pdf>.
- Sutton, R.S. and Barto, A. G. *Reinforcement Learning: An Introduction*. MIT Press, 1998.
- Thompson, W. R. On the likelihood that one unknown probability exceeds another in view of the evidence of two samples. *Biometrika*, 25:285–294, 1933.
- Wilson, Aaron, Fern, Alan, Ray, Soumya, and Tadepalli, Prasad. Multi-task reinforcement learning: a hierarchical bayesian approach. In *ICML '07: Proceedings of the 24th international conference on Machine learning*, pp. 1015–1022, New York, NY, USA, 2007. ACM. ISBN 978-1-59593-793-3. doi: <http://doi.acm.org/10.1145/1273496.1273624>.